

1 Start coding or generate with AI.

```
1 # 1. Fresh install kaggle
2 !pip uninstall -y kaggle
3 !pip install --upgrade kaggle
4
5 # 2. Upload kaggle.json
6 from google.colab import files
7 files.upload() # upload kaggle.json
8
9 !mkdir -p ~/.kaggle
10 !cp kaggle.json ~/.kaggle/
11 !chmod 600 ~/.kaggle/kaggle.json
12
13 # 3. Download & unzip new dataset
14 !rm -rf TB_Chest_Xray
15 !kaggle datasets download -d tawsifurrahman/tuberculosis-tb-chest-xray-dataset
16 !unzip -q tuberculosis-tb-chest-xray-dataset.zip -d ./TB_Chest_Xray
17
18 # 4. Check folder structure
19 !ls ./TB_Chest_Xray/Tuberculosis_Chest_Radiography_Database
20
```

```
1 import os, shutil
2 from sklearn.model_selection import train_test_split
3
4 data_dir = "./TB_Chest_Xray/TB_Chest_Radiography_Database"
5 split_dir = "./TB_SPLIT"
6 os.makedirs(split_dir, exist_ok=True)
7
8 for cls in ["Normal", "Tuberculosis"]:
9     cls_dir = os.path.join(data_dir, cls)
10    images = [os.path.join(cls_dir, f) for f in os.listdir(cls_dir) if f.lower().endswith(".png", ".JPG")]
11
12    # First split: train vs temp (val+test)
13    train, temp = train_test_split(images, test_size=0.2, random_state=42, stratify=[cls]*len(images))
14
15    # Second split: val vs test
16    val, test = train_test_split(temp, test_size=0.5, random_state=42, stratify=[cls]*len(temp))
```

```

16
17     for split, split_files in zip(["train", "val", "test"], [train, val, test]):
18         split_path = os.path.join(split_dir, split, cls)
19         os.makedirs(split_path, exist_ok=True)
20         for img in split_files:
21             shutil.copy(img, split_path)
22
23 print("Dataset split created at:", split_dir)
24

```

Dataset split created at: ./TB_SPLIT

```

1 import torch
2 from torch.utils.data import DataLoader
3 import torchvision.transforms as transforms
4 import torchvision.datasets as datasets
5
6 # Transforms
7 transform_train = transforms.Compose([
8     transforms.Resize((224,224)),
9     transforms.RandomHorizontalFlip(),
10    transforms.ToTensor(),
11    transforms.Normalize(mean=[0.485, 0.456, 0.406],
12                         std=[0.229, 0.224, 0.225]),
13 ])
14
15 transform_test = transforms.Compose([
16     transforms.Resize((224,224)),
17     transforms.ToTensor(),
18     transforms.Normalize(mean=[0.485, 0.456, 0.406],
19                         std=[0.229, 0.224, 0.225]),
20 ])
21
22 # Datasets (ImageFolder picks class names from subfolders)
23 trainset = datasets.ImageFolder(os.path.join(split_dir, "train"), transform=transform_train)
24 valset   = datasets.ImageFolder(os.path.join(split_dir, "val"), transform=transform_test)
25 testset  = datasets.ImageFolder(os.path.join(split_dir, "test"), transform=transform_test)
26
27 # Dataloaders
28 trainloader = DataLoader(trainset, batch_size=32, shuffle=True, num_workers=2)
29 valloader   = DataLoader(valset, batch_size=32, shuffle=False, num_workers=2)
30 testloader  = DataLoader(testset, batch_size=32, shuffle=False, num_workers=2)
31
32 print("Train size:", len(trainset))
33 print("Val size:", len(valset))
34 print("Test size:", len(testset))
35 print("Classes:", trainset.classes)
36

```

Train size: 3360
 Val size: 420
 Test size: 420
 Classes: ['Normal', 'Tuberculosis']

```

1 # =====
2 # 4. CvT Model Definition
3 # =====
4 import torch
5 import torch.nn as nn
6 import torch.nn.functional as F
7 import torch.optim as optim
8 from collections import Counter
9
10 # This is a dummy class to represent your custom CvT model.
11 # The user's provided code is copied here with the fix in the forward pass.
12 class CvT(nn.Module):
13     def __init__(self, num_classes=2):

```

```

14     super().__init__()
15     # Simplified for demonstration. This is where your CvT layers go.
16     # Your provided CvT model definition is integrated below.
17     self.embed = ConvTokenEmbed(3, 64, kernel_size=7, stride=4)
18     self.blocks = nn.Sequential(
19         CvTBlock(64, num_heads=2),
20         CvTBlock(64, num_heads=2),
21     )
22     self.norm = nn.LayerNorm(64)
23     self.cls_token = nn.Parameter(torch.zeros(1, 1, 64))
24     self.fc = nn.Linear(64, num_classes)
25
26     def forward(self, x):
27         B = x.shape[0]
28         x = self.embed(x) # [B, N, 64] where N is number of image tokens
29
30         # --- FIX: Correctly separating and re-concatenating tokens ---
31         # Acknowledge and revert the previous bug introduced in the forward pass.
32         # The class token (cls_tokens) and image tokens (x) must be handled separately
33         # before passing through the attention blocks to avoid the shape error.
34         # The ConvAttention module expects a perfect square number of tokens.
35
36         cls_tokens = self.cls_token.expand(B, -1, -1)
37
38         # Pass only the image tokens to the CvTBlocks, which use ConvAttention.
39         img_tokens_out = self.blocks(x)
40
41         # Now, concatenate the class token back with the processed image tokens.
42         x = torch.cat((cls_tokens, img_tokens_out), dim=1)
43
44         x = self.norm(x)
45
46         # The final classification is based on the class token's output.
47         return self.fc(x[:, 0])
48
49 class ConvTokenEmbed(nn.Module):
50     def __init__(self, in_channels, embed_dim, kernel_size=7, stride=4, padding=3):
51         super().__init__()
52         self.conv = nn.Conv2d(in_channels, embed_dim, kernel_size, stride, padding, bias=False)
53         self.bn = nn.BatchNorm2d(embed_dim)
54     def forward(self, x):
55         x = F.relu(self.bn(self.conv(x)))
56         return x.flatten(2).transpose(1, 2)
57
58 class ConvAttention(nn.Module):
59     def __init__(self, embed_dim, num_heads):
60         super().__init__()
61         self.embed_dim = embed_dim
62         self.num_heads = num_heads
63         self.scale = (embed_dim // num_heads) ** -0.5
64         self.depthwise_conv = nn.Conv2d(embed_dim, embed_dim, kernel_size=3, padding=1, groups=embed_dim)
65         self.pointwise_conv = nn.Conv2d(embed_dim, embed_dim * 3, kernel_size=1)
66         self.proj = nn.Linear(embed_dim, embed_dim)
67     def forward(self, x):
68         B, N, C = x.shape
69         x_reshaped = x.transpose(1, 2).view(B, C, int(N ** 0.5), int(N ** 0.5))
70         qkv = self.pointwise_conv(self.depthwise_conv(x_reshaped)).view(B, 3 * C, N).transpose(1, 2)
71         qkv = qkv.reshape(B, N, 3, self.num_heads, C // self.num_heads).permute(2, 0, 3, 1, 4)
72         q, k, v = qkv[0], qkv[1], qkv[2]
73         attn = (q @ k.transpose(-2, -1)) * self.scale
74         attn = F.softmax(attn, dim=-1)
75         x = (attn @ v).transpose(1, 2).reshape(B, N, C)
76         return self.proj(x)
77
78 class CvTBlock(nn.Module):
79     def __init__(self, embed_dim, num_heads):

```

```

80     super().__init__()
81     self.norm1 = nn.LayerNorm(embed_dim)
82     self.attn = ConvAttention(embed_dim, num_heads)
83     self.norm2 = nn.LayerNorm(embed_dim)
84     self.mlp = nn.Sequential(
85         nn.Linear(embed_dim, 4 * embed_dim),
86         nn.GELU(),
87         nn.Linear(4 * embed_dim, embed_dim)
88     )
89     def forward(self, x):
90         x = x + self.attn(self.norm1(x))
91         x = x + self.mlp(self.norm2(x))
92         return x
93
94
95 # =====
96 # 5. Training Setup
97 # =====
98 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
99 model = CvT(num_classes=2).to(device)
100
101 # --- FIX: Weighted Loss Function to address data imbalance ---
102 # Assuming a 5:1 ratio of majority to minority class.
103 # We'll assign a weight of 5 to the minority class (index 1) and 1 to the majority (index 0).
104 # You should replace these values with the actual inverse class frequencies from your dataset.
105 weights = torch.tensor([1.0, 5.0]).to(device)
106 criterion = nn.CrossEntropyLoss(weight=weights)
107
108
109 optimizer = optim.SGD(model.parameters(), lr=0.05, momentum=0.9, weight_decay=5e-3)
110 scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=20, gamma=0.1)
111
112 def train_one_epoch(model, loader, criterion, optimizer, device):
113     model.train()
114     running_loss, correct, total = 0.0, 0, 0
115     for images, labels in loader:
116         images, labels = images.to(device), labels.to(device)
117         optimizer.zero_grad()
118         outputs = model(images)
119         loss = criterion(outputs, labels)
120         loss.backward()
121         optimizer.step()
122         running_loss += loss.item()
123         _, predicted = torch.max(outputs, 1)
124         total += labels.size(0)
125         correct += (predicted == labels).sum().item()
126     return running_loss / len(loader), 100 * correct / total
127
128 def evaluate(model, loader, criterion, device):
129     model.eval()
130     running_loss, correct, total = 0.0, 0, 0
131     with torch.no_grad():
132         for images, labels in loader:
133             images, labels = images.to(device), labels.to(device)
134             outputs = model(images)
135             loss = criterion(outputs, labels)
136             running_loss += loss.item()
137             _, predicted = torch.max(outputs, 1)
138             total += labels.size(0)
139             correct += (predicted == labels).sum().item()
140     return running_loss / len(loader), 100 * correct / total
141

```

```

1 import matplotlib.pyplot as plt
2 import numpy as np
3 import seaborn as sns

```

```
4 import torch
5
6 # =====
7 # 1. Class Distribution
8 # =====
9 def plot_class_distribution(trainset, valset, testset):
10    def get_labels(dataset):
11        return [label for _, label in dataset]
12
13    datasets_map = {
14        "Train": get_labels(trainset),
15        "Validation": get_labels(valset),
16        "Test": get_labels(testset),
17    }
18
19    plt.figure(figsize=(15, 4))
20    for i, (name, labels) in enumerate(datasets_map.items()):
21        plt.subplot(1, 3, i + 1)
22        sns.countplot(x=labels, palette="viridis")
23        plt.title(f"{name} Set Distribution")
24        plt.xticks(
25            ticks=range(len(trainset.classes)),
26            labels=trainset.classes,
27            rotation=20
28        )
29    plt.tight_layout()
30    plt.show()
31
32 plot_class_distribution(trainset, valset, testset)
33
34
35 # =====
36 # 2. Sample Images
37 # =====
38 def show_sample_images(dataset, class_names, n=5):
39    plt.figure(figsize=(15, 3))
40    for i in range(n):
41        img, label = dataset[i]
42        img = img.permute(1, 2, 0).numpy()
43        img = np.clip(img * 0.229 + 0.485, 0, 1) # unnormalize (approximate)
44        plt.subplot(1, n, i + 1)
45        plt.imshow(img)
46        plt.title(class_names[label])
47        plt.axis("off")
48    plt.show()
49
50 show_sample_images(trainset, trainset.classes)
51
52
53 # =====
54 # 3. Intensity Histograms
55 # =====
56 def plot_intensity_histograms(dataset, n_samples=200):
57    imgs = [dataset[i][0].permute(1, 2, 0).numpy() for i in range(min(n_samples, len(dataset)))]
58    pixels = np.vstack([img.reshape(-1, 3) for img in imgs])
59
60    plt.figure(figsize=(12, 4))
61    colors = ["r", "g", "b"]
62    for i, c in enumerate(colors):
63        plt.hist(pixels[:, i], bins=50, color=c, alpha=0.5, label=f"{c}-channel")
64    plt.title("Pixel Intensity Distribution (sample)")
65    plt.legend()
66    plt.show()
67
68 plot_intensity_histograms(trainset)
69
```

```
70
71 # =====
72 # 4. Class Balance Heatmap
73 # =====
74 def plot_class_heatmap(trainset, valset, testset):
75     def count_labels(dataset):
76         counts = torch.bincount(torch.tensor([y for _, y in dataset]))
77         return counts.numpy()
78
79     data = np.vstack([
80         count_labels(trainset),
81         count_labels(valset),
82         count_labels(testset)
83     ])
84
85     sns.heatmap(data, annot=True, fmt="d", cmap="Blues",
86                 xticklabels=trainset.classes,
87                 yticklabels=["Train", "Val", "Test"])
88     plt.title("Class Distribution Heatmap")
89     plt.show()
90
91 plot_class_heatmap(trainset, valset, testset)
92
93
94 # =====
95 # 5. Dataset Size Summary
96 # =====
97 print(f"Train set size: {len(trainset)}")
98 print(f"Validation set size: {len(valset)}")
99 print(f"Test set size: {len(testset)}")
100 print(f"Classes: {trainset.classes}")
101
```



```
/tmp/ipython-input-4141679274.py:22: FutureWarning:  
  Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set  
  sns.countplot(x=labels, palette="viridis")  
/tmp/ipython-input-4141679274.py:22: FutureWarning:  
  Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set  
  sns.countplot(x=labels, palette="viridis")  
/tmp/ipython-input-4141679274.py:22: FutureWarning:  
  Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set  
  sns.countplot(y=labels, palette="viridis")  
1 from sklearn.metrics import confusion_matrix, classification_report, roc_curve, auc  
2 import seaborn as sns  
3 import numpy as np  
4 import matplotlib.pyplot as plt  
5 import torch  
6  
7 # =====  
8 # 1. Training with logging  
9 # =====  
10 num_epochs = 20  
11 history = {  
12     "train_loss": [], "train_acc": [],  
13     "val_loss": [], "val_acc": [],  
14     "test_loss": [], "test_acc": []  
15 }  
16  
17 for epoch in range(1, num_epochs + 1):  
18     train_loss, train_acc = train_one_epoch(model, trainloader, criterion, optimizer, device)  
19     val_loss, val_acc = evaluate(model, valloader, criterion, device)  
20     test_loss, test_acc = evaluate(model, testloader, criterion, device)  
21  
22     scheduler.step()  
23  
24     # log  
25     history["train_loss"].append(train_loss)  
26     history["train_acc"].append(train_acc)  
27     history["val_loss"].append(val_loss)  
28     history["val_acc"].append(val_acc)  
29     history["test_loss"].append(test_loss)  
30     history["test_acc"].append(test_acc)  
31  
32     print(f"Epoch {epoch:02d}: "  
33           f"Train Acc {train_acc:.2f}% | "  
34           f"Val Acc {val_acc:.2f}% | "  
35           f"Test Acc {test_acc:.2f}%")  
36  
37  
38 # =====  
39 # 2. Loss & Accuracy Curves  
40 # =====  
41 def plot_training_curves(history):  
42     epochs = range(1, len(history["train_loss"]) + 1)  
43  
44     plt.figure(figsize=(12,5))  
45     plt.subplot(1,2,1)  
46     plt.plot(epochs, history["train_loss"], label="Train")  
47     plt.plot(epochs, history["val_loss"], label="Val")  
48     plt.plot(epochs, history["test_loss"], label="Test")  
49     plt.title("Loss")  
50     plt.xlabel("Epoch")  
51     plt.ylabel("Loss")  
52     plt.legend()  
53  
54     plt.subplot(1,2,2)  
55     plt.plot(epochs, history["train_acc"], label="Train")
```

```

55 plt.plot(epochs, history["train_acc"], label="Train")
56 plt.plot(epochs, history["val_acc"], label="Val")
57 plt.plot(epochs, history["test_acc"], label="Test")
58 plt.title("Accuracy")
59 plt.xlabel("Epoch")
60 plt.ylabel("Accuracy (%)")
61 plt.legend()
62
63 plt.show()
64
65 plot_training_curves(history)
66
67
68 # =====
69 # 3. Confusion Matrix & Report
70 # =====
71 def evaluate_predictions(model, loader, device, class_names):
72     model.eval()
73     all_preds, all_labels = [], []
74     with torch.no_grad():
75         for images, labels in loader:
76             images, labels = images.to(device), labels.to(device)
77             outputs = model(images)
78             preds = outputs.argmax(dim=1)
79             all_preds.extend(preds.cpu().numpy())
80             all_labels.extend(labels.cpu().numpy())
81     return np.array(all_preds), np.array(all_labels)
82
83 # Use test set for evaluation
84 y_pred, y_true = evaluate_predictions(model, testloader, device, trainset.classes)
85
86 # Confusion Matrix
87 cm = confusion_matrix(y_true, y_pred)
88 plt.figure(figsize=(6,5))
89 sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
90             xticklabels=trainset.classes, yticklabels=trainset.classes)
91 plt.xlabel("Predicted")
92 plt.ylabel("True")
93 plt.title("Confusion Matrix")
94 plt.show()
95
96 # Classification Report
97 print("Classification Report:\n")
98 print(classification_report(y_true, y_pred, target_names=trainset.classes))
99
100
101 # =====
102 # 4. ROC Curve (binary)
103 # =====
104 if len(trainset.classes) == 2:
105     from sklearn.preprocessing import label_binarize
106
107     y_true_bin = label_binarize(y_true, classes=[0,1]).ravel()
108     y_score = []
109     model.eval()
110     with torch.no_grad():
111         for images, labels in testloader:
112             images = images.to(device)
113             outputs = model(images)
114             probs = torch.softmax(outputs, dim=1)[:,1] # probability of class 1
115             y_score.extend(probs.cpu().numpy())
116     y_score = np.array(y_score)
117
118     fpr, tpr, _ = roc_curve(y_true_bin, y_score)
119     roc_auc = auc(fpr, tpr)
120
121     plt.figure(figsize=(6,5))

```

```
122 plt.plot(fpr, tpr, label="ROC Curve (AUC = {:.2f})")
123 plt.plot([0,1], [0,1], linestyle="--", color="gray")
124 plt.xlabel("False Positive Rate")
125 plt.ylabel("True Positive Rate")
126 plt.title("ROC Curve")
127 plt.legend()
128 plt.show()
129
130
131 # =====
132 # 5. Misclassified Samples
133 # =====
134 def show_misclassified(model, loader, class_names, device, n=6):
135     model.eval()
136     misclassified = []
137     with torch.no_grad():
138         for images, labels in loader:
139             images, labels = images.to(device), labels.to(device)
140             outputs = model(images)
141             preds = outputs.argmax(dim=1)
142             for img, pred, true in zip(images, preds, labels):
143                 if pred != true:
144                     misclassified.append((img.cpu(), pred.item(), true.item()))
145             if len(misclassified) >= n:
146                 break
147             if len(misclassified) >= n:
148                 break
149
150     plt.figure(figsize=(15, 5))
151     for i, (img, pred, true) in enumerate(misclassified):
152         plt.subplot(1, n, i+1)
153         img = img.permute(1,2,0).numpy()
154         img = np.clip(img*0.229 + 0.485, 0, 1) # unnormalize approx
155         plt.imshow(img)
156         plt.title(f"P:{class_names[pred]}\nT:{class_names[true]}")
157         plt.axis("off")
158     plt.suptitle("Misclassified Samples", fontsize=16)
159     plt.show()
160
161 show_misclassified(model, testloader, trainset.classes, device, n=6)
162
163
```